




# Exercise 1

# Goal of Machine Learning

- ▶ 다음 데이터를 가장 잘 설명하는 함수를 찾아라

(0.0, 0.0) (1.0, 1.0) (1.0, 2.0)



$$f(x; w_0, w_1) = w_1 x + w_0$$

$$Error = \sum_{(x,y) \in Data} (y - f(\mathbf{x}; w_1, w_2, \dots, w_m))^2$$

# Gradient Descent Method

## ► Steps

Randomly choose an initial solution,  $w^0$

Repeat

$$w^{t+1} = w^t - \eta \left. \frac{dE}{dw} \right|_{w=w^t}$$


learning rate

Until **stopping condition** is satisfied

- $|w^{t+1} - w^t|$  is very small
- $f(w)$  little moves
- fixed number of iterations

# Coding 전에 생각해 볼 것

- ▶ 입력 데이터 : (0.0), (1.0), (1.0)
- ▶ 출력 데이터 : (0.0), (1.0), (2.0)
- ▶ Optimizer: gradient descent method
- ▶ Loss function: Mean square error



# 준비 단계 1

▶ 입력 데이터 : (0.0), (1.0), (1.0)

▶ 출력 데이터 : (0.0), (1.0), (2.0)

```
x_train = torch.FloatTensor([[0], [1], [1]])  
y_train = torch.FloatTensor([[0], [1], [2]])
```

▶ Model:  $f(X; W_0, W_1) = W_1 * X + W_0$

```
hypothesis = x_train * w + b
```

▶ Loss function :  $\text{Error} = \sum (y - f(X; W_1, W_2, \dots, W_m))^2$

```
cost = torch.mean( (hypothesis - y_train)**2)
```

# optimizer

## ▶ Optimizer : SGD

### Algorithms

Adadelta	Implements Adadelta algorithm.
Adagrad	Implements Adagrad algorithm.
Adam	Implements Adam algorithm.
AdamW	Implements AdamW algorithm.
SparseAdam	Implements lazy version of Adam algorithm suitable for sparse tensors.
Adamax	Implements Adamax algorithm (a variant of Adam based on infinity norm).
ASGD	Implements Averaged Stochastic Gradient Descent.
LBFGS	Implements L-BFGS algorithm, heavily inspired by <code>minFunc</code> .
RMSprop	Implements RMSprop algorithm.
Rprop	Implements the resilient backpropagation algorithm.
SGD	Implements stochastic gradient descent (optionally with momentum).

<https://pytorch.org/docs/stable/optim.html?highlight=optimizer#torch.optim.Optimizer>

# optimizer

## ▶ Optimizer : SGD

Docs > torch.optim > SGD



### SGD

**CLASS** `torch.optim.SGD(params, lr=<required parameter>, momentum=0, dampening=0, weight_decay=0, nesterov=False)`

[SOURCE]

Implements stochastic gradient descent (optionally with momentum).

Nesterov momentum is based on the formula from [On the importance of initialization and momentum in deep learning](#).

#### Parameters

- **params** (*iterable*) – iterable of parameters to optimize or dicts defining parameter groups
- **lr** (*float*) – learning rate
- **momentum** (*float, optional*) – momentum factor (default: 0)
- **weight\_decay** (*float, optional*) – weight decay (L2 penalty) (default: 0)
- **dampening** (*float, optional*) – dampening for momentum (default: 0)
- **nesterov** (*bool, optional*) – enables Nesterov momentum (default: False)

#### Example

```
>>> optimizer = torch.optim.SGD(model.parameters(), lr=0.1, momentum=0.9)
>>> optimizer.zero_grad()
>>> loss_fn(model(input), target).backward()
>>> optimizer.step()
```

# Exercise 1

```
import torch
import numpy as np
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

x_train = torch.FloatTensor([[0], [1], [1]])
y_train = torch.FloatTensor([[0], [1], [2]])

w = torch.zeros(1, requires_grad = True)
b = torch.zeros(1, requires_grad = True)

optimizer = optim.SGD([w, b] , lr = 0.01)

nb_epochs = 1000

for epoch in range( nb_epochs + 1):

    hypothesis = x_train * w + b

    cost = torch.mean( (hypothesis - y_train)**2)

    optimizer.zero_grad()
    cost.backward()
    optimizer.step()

    if epoch % 100 == 0 :
        print( 'Epoch {:5d}/{:} w:{:.3f} b:{:.3f} cost: {:.3f}'
              .format(epoch, nb_epochs, w.item(), b.item(), cost.item() ))
```

$$f(x; w_0, w_1) = w_1 x + w_0$$



# Exercise 1

---

Epoch	0/3000	w:0.020	b:0.020	cost: 1.667
Epoch	100/3000	w:0.782	b:0.509	cost: 0.283
Epoch	200/3000	w:0.981	b:0.403	cost: 0.230
Epoch	300/3000	w:1.114	b:0.301	cost: 0.202
Epoch	400/3000	w:1.212	b:0.225	cost: 0.186
Epoch	500/3000	w:1.285	b:0.168	cost: 0.178
Epoch	600/3000	w:1.340	b:0.125	cost: 0.173
Epoch	700/3000	w:1.380	b:0.093	cost: 0.170
Epoch	800/3000	w:1.411	b:0.070	cost: 0.169
Epoch	900/3000	w:1.433	b:0.052	cost: 0.168
Epoch	1000/3000	w:1.450	b:0.039	cost: 0.167
Epoch	1100/3000	w:1.463	b:0.029	cost: 0.167
Epoch	1200/3000	w:1.472	b:0.022	cost: 0.167
Epoch	1300/3000	w:1.479	b:0.016	cost: 0.167
Epoch	1400/3000	w:1.485	b:0.012	cost: 0.167
Epoch	1500/3000	w:1.488	b:0.009	cost: 0.167
Epoch	1600/3000	w:1.491	b:0.007	cost: 0.167
Epoch	1700/3000	w:1.494	b:0.005	cost: 0.167
Epoch	1800/3000	w:1.495	b:0.004	cost: 0.167
Epoch	1900/3000	w:1.496	b:0.003	cost: 0.167
Epoch	2000/3000	w:1.497	b:0.002	cost: 0.167
Epoch	2100/3000	w:1.498	b:0.002	cost: 0.167
Epoch	2200/3000	w:1.499	b:0.001	cost: 0.167
Epoch	2300/3000	w:1.499	b:0.001	cost: 0.167
Epoch	2400/3000	w:1.499	b:0.001	cost: 0.167
Epoch	2500/3000	w:1.499	b:0.000	cost: 0.167
Epoch	2600/3000	w:1.500	b:0.000	cost: 0.167
Epoch	2700/3000	w:1.500	b:0.000	cost: 0.167
Epoch	2800/3000	w:1.500	b:0.000	cost: 0.167
Epoch	2900/3000	w:1.500	b:0.000	cost: 0.167
Epoch	3000/3000	w:1.500	b:0.000	cost: 0.167

# Quiz 1

## ► Data : (x1, x2, x3) (Y)

```
x_train = torch.FloatTensor([[73,80,75], [93,88,93], [89, 91, 90], [96,98,100], [73,66,70]])  
y_train = torch.FloatTensor([[152], [185], [180], [196], [142]])
```

## ► hint

- $[a1, a2, a3] * [w1, w2, w3] = ?$
- MSELOSS ?

Docs > torch.nn > MSELoss



## MSELOSS

**CLASS** `torch.nn.MSELoss(size_average=None, reduce=None, reduction='mean')`

[\[SOURCE\]](#)

Creates a criterion that measures the mean squared error (squared L2 norm) between each element in the input  $x$  and target  $y$ .

The unreduced (i.e. with `reduction` set to `'none'`) loss can be described as:

$$\ell(x, y) = L = \{l_1, \dots, l_N\}^\top, \quad l_n = (x_n - y_n)^2,$$

where  $N$  is the batch size. If `reduction` is not `'none'` (default `'mean'`), then:

$$\ell(x, y) = \begin{cases} \text{mean}(L), & \text{if reduction = 'mean'}; \\ \text{sum}(L), & \text{if reduction = 'sum'}. \end{cases}$$

Examples:

```
>>> loss = nn.MSELoss()  
>>> input = torch.randn(3, 5, requires_grad=True)  
>>> target = torch.randn(3, 5)  
>>> output = loss(input, target)  
>>> output.backward()
```

# Answer 1

```
import torch
import numpy as np
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

x_train = torch.FloatTensor([[73,80,75], [93,88,93], [89, 91, 90], [96,98,100], [73,66,70]])
y_train = torch.FloatTensor([[152], [185], [180], [196], [142]])

w = torch.zeros((3, 1), requires_grad = True)
b = torch.zeros(1, requires_grad = True)

optimizer = optim.SGD([w, b], lr = 1e-5)

nb_epochs = 1000

for epoch in range( nb_epochs + 1):

    #hypothesis = x_train * w + b
    hypothesis = x_train.matmul(w) + b

    #cost = torch.mean( (hypothesis - y_train)**2)
    cost = F.mse_loss(hypothesis, y_train)

    optimizer.zero_grad()
    cost.backward()
    optimizer.step()

    if epoch % 100 == 0 :
        print( 'Epoch {:5d}/{:} cost: {:.3f}'
              .format(epoch, nb_epochs, cost.item() ))
        print( w.squeeze() )
```

```
loss = nn.MSELoss()

nb_epochs = 1000

for epoch in range( nb_epochs + 1):

    #hypothesis = x_train * w + b
    hypothesis = x_train.matmul(w) + b

    #cost = torch.mean( (hypothesis - y_train)**2)
    #cost = F.mse_loss(hypothesis, y_train)
    cost = loss(hypothesis, y_train)
```

# MSE Loss

- ▶ **Torch.nn.function : 함수, torch.nn : 클래스**

## SOURCE CODE FOR TORCH.NN.MODULES.LOSS

```
import warnings

from .distance import PairwiseDistance
from .module import Module
from .. import functional as F
from .. import _reduction as _Reduction
```

```
class MSELoss(_loss):
    """Creates a criterion that measures the mean squared error (squared L2 norm) between
    each element in the input :math:`x` and target :math:`y`.
```

```
    __constants__ = ['reduction']

    def __init__(self, size_average=None, reduce=None, reduction: str = 'mean') -> None:
        super(MSELoss, self).__init__(size_average, reduce, reduction)

    def forward(self, input: Tensor, target: Tensor) -> Tensor:
        return F.mse_loss(input, target, reduction=self.reduction)
```

# SOURCE CODE FOR TORCH.NN.FUNCTIONAL

```
r"""Functional interface"""
from typing import Callable, List, Optional, Tuple
import math
import warnings
```

```
def mse_loss(
    input: Tensor,
    target: Tensor,
    size_average: Optional[bool] = None,
    reduce: Optional[bool] = None,
    reduction: str = "mean",
) -> Tensor:
    r"""mse_loss(input, target, size_average=None, reduce=None, reduction='mean') -> Tensor

    Measures the element-wise mean squared error.

    See :class:`~torch.nn.MSELoss` for details.
    """
    if has_torch_function_variadic(input, target):
        return handle_torch_function(
            mse_loss, (input, target), input, target, size_average=size_average, reduce=reduce,
            reduction=reduction
        )
    if not (target.size() == input.size()):
        warnings.warn(
            "Using a target size ({}) that is different to the input size ({}). ".
            "This will likely lead to incorrect results due to broadcasting. ".
            "Please ensure they have the same size.".format(target.size(), input.size()),
            stacklevel=2,
        )
    if size_average is not None or reduce is not None:
        reduction = _Reduction.legacy_get_string(size_average, reduce)

    expanded_input, expanded_target = torch.broadcast_tensors(input, target)
    return torch._C._nn.mse_loss(expanded_input, expanded_target,
        _Reduction.get_enum(reduction))
```

# Tip

---

## ▶ Data Preprocessing


```
# 데이터
x_train = torch.FloatTensor([[73, 80, 75],
                             [93, 88, 93],
                             [89, 91, 90],
                             [96, 98, 100],
                             [73, 66, 70]])

mu = x_train.mean(dim=0)
sigma = x_train.std(dim=0)
norm_x_train = (x_train - mu) / sigma
print(norm_x_train)
```

# PyTorch Module

## ▶ How can we make Deep Neural Network?

```
W = torch.zeros((3, 1), requires_grad=True)  
b = torch.zeros(1, requires_grad=True)
```



# PyTorch Module

---

## ▶ nn.Module (1)

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

class MultivariateLinearRegressionModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.linear = nn.Linear(3, 1)

    def forward(self, x):
        return self.linear(x)
```



# PyTorch Module

## ▶ nn.Module (2)

```
# 데이터
x_train = torch.FloatTensor([[73, 80, 75],
                              [93, 88, 93],
                              [89, 91, 90],
                              [96, 98, 100],
                              [73, 66, 70]])
y_train = torch.FloatTensor([[152], [185], [180], [196], [142]])
# 모델 초기화
model = MultivariateLinearRegressionModel()
# optimizer 설정
optimizer = optim.SGD(model.parameters(), lr=1e-5)
```

# PyTorch Module

## ▶ nn.Module (3)

```
nb_epochs = 20
for epoch in range(nb_epochs + 1):
    # H(x) 계산
    prediction = model(x_train)

    # cost 계산
    cost = F.mse_loss(prediction, y_train)

    # cost로 H(x) 개선
    optimizer.zero_grad()
    cost.backward()
    optimizer.step()

    # 20번마다 로그 출력
    print('Epoch {:4d}/{:4d} Cost: {:.6f}'.format(
        epoch, nb_epochs, cost.item()
    ))
```

# Pytorch

---

## ▶ GPU

```
import torch
import numpy as np

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

print(device)

#net = Model.to(device)
#inputs = data.to(device)
```

# Question and Answer